

simpledb

环境配置

使用Idea

使用Eclipse

介绍

Exercise 1

Exercise 2

Exercise 3

Exercise 4

Exercise 5

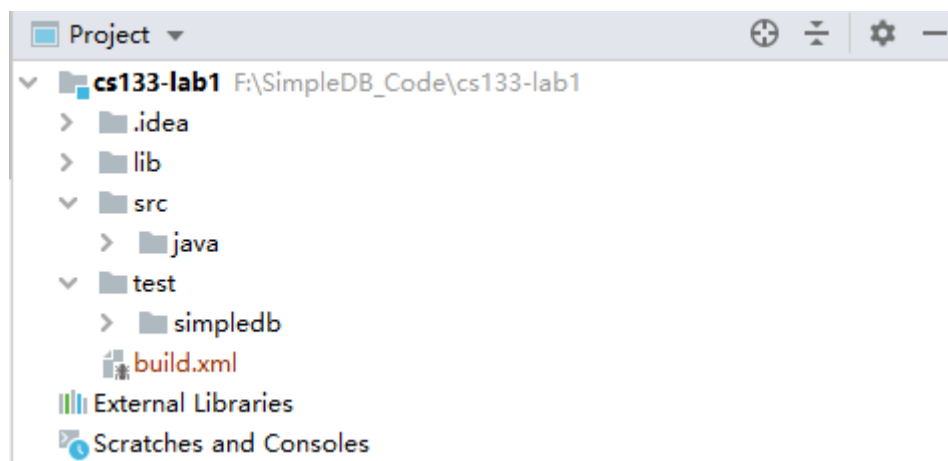
Exercise 6

simpledb

环境配置

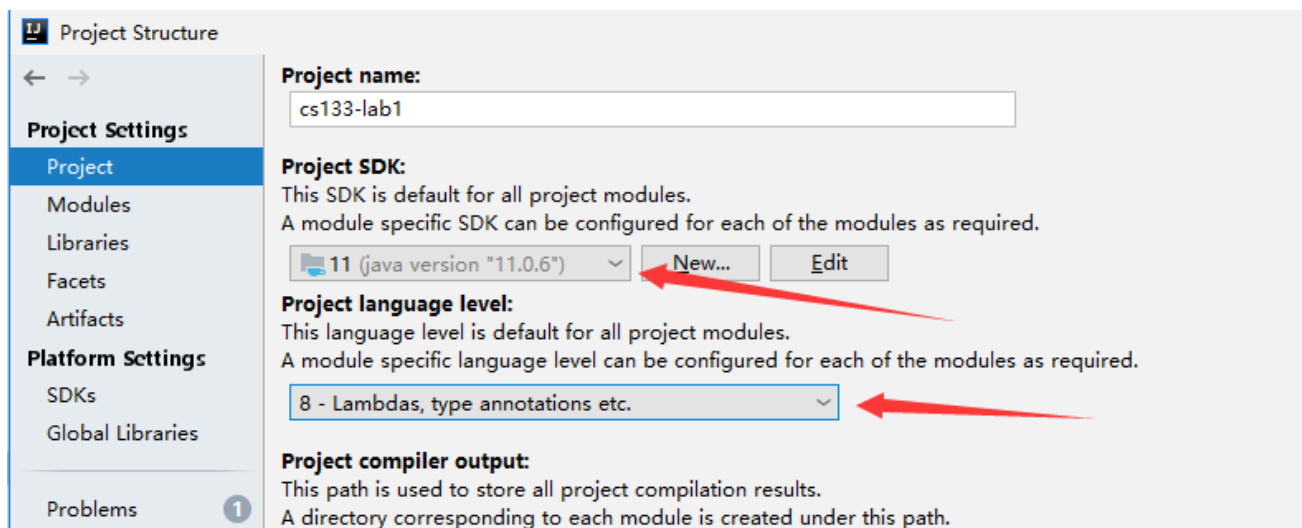
使用Idea

打开工程

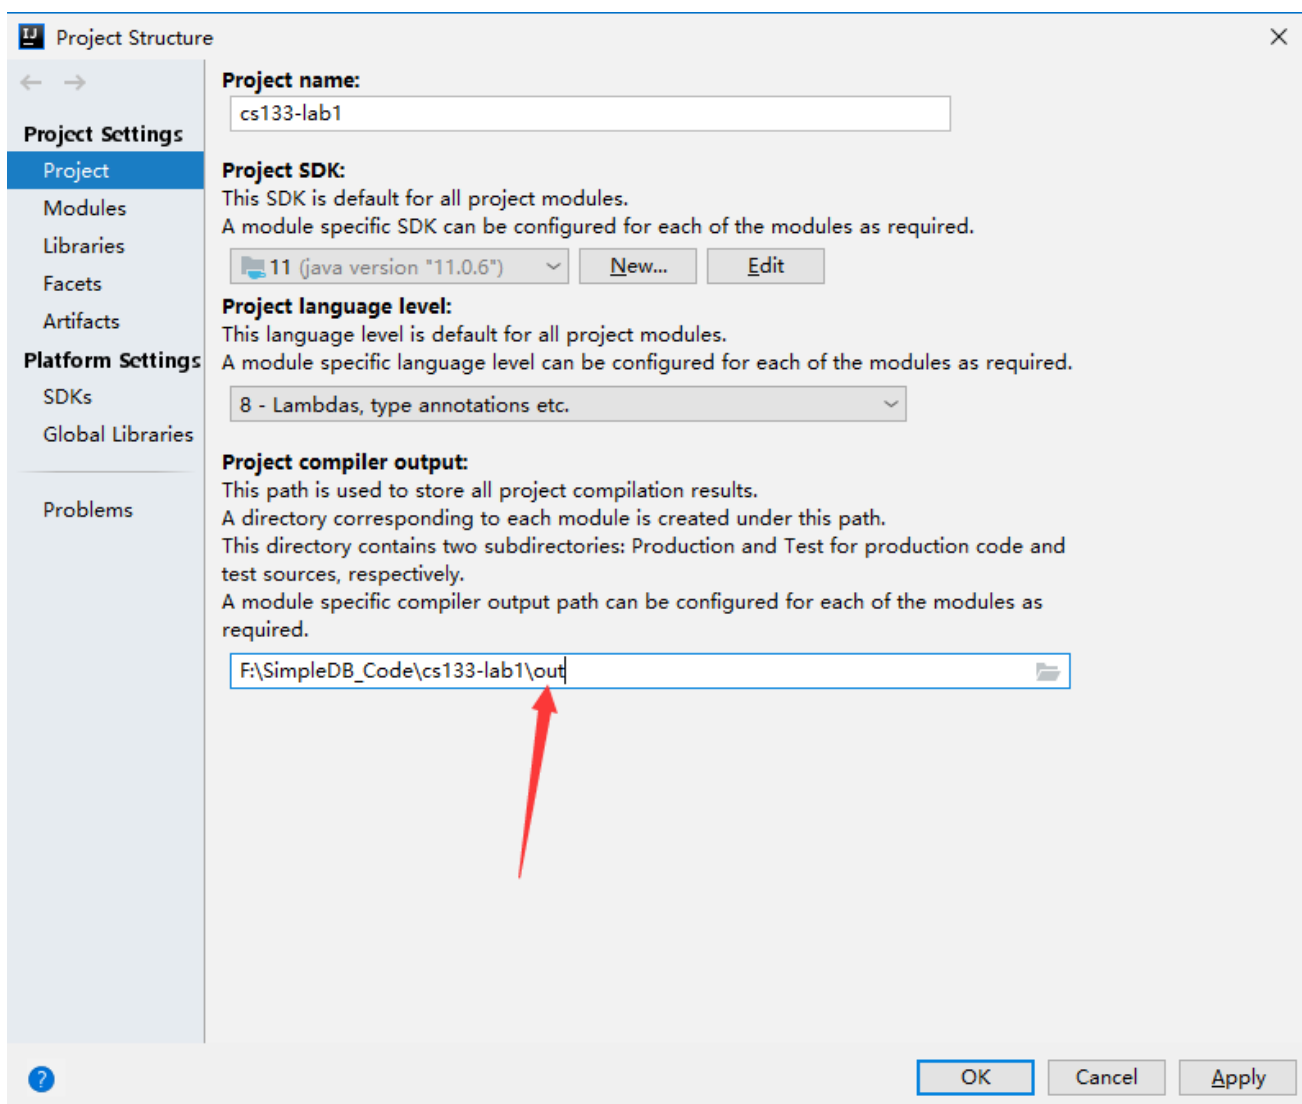


ctrl+alt+shift+s 打开 Project Structure

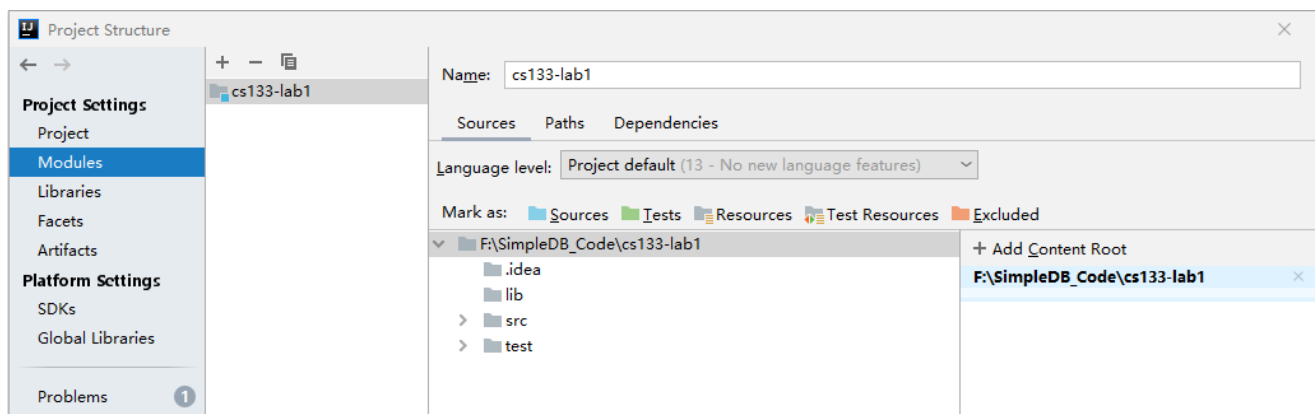
配置 SDK



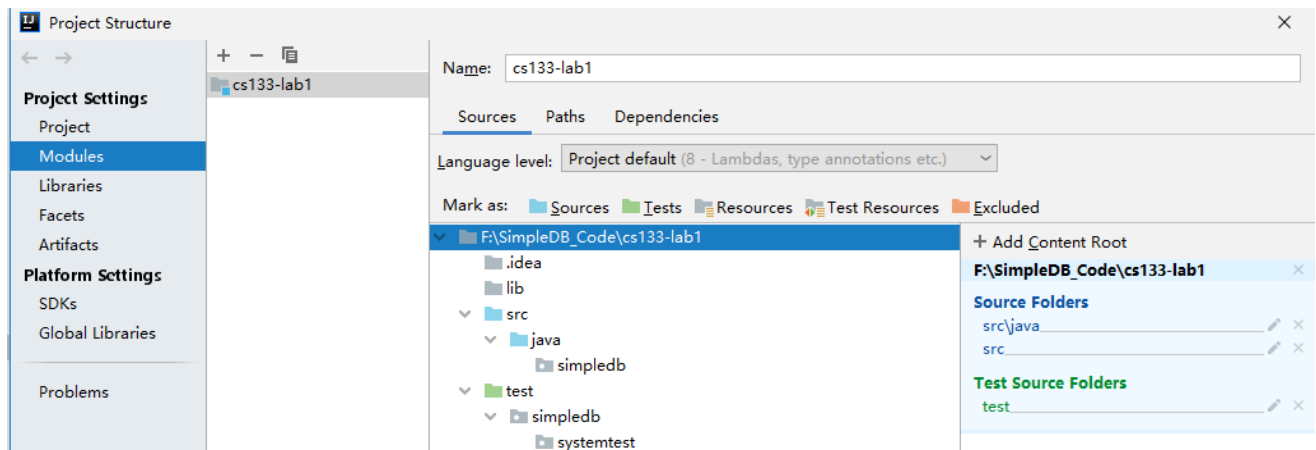
然后配置 Project compiler output，路径我们配置成 当前项目路径\out 即可



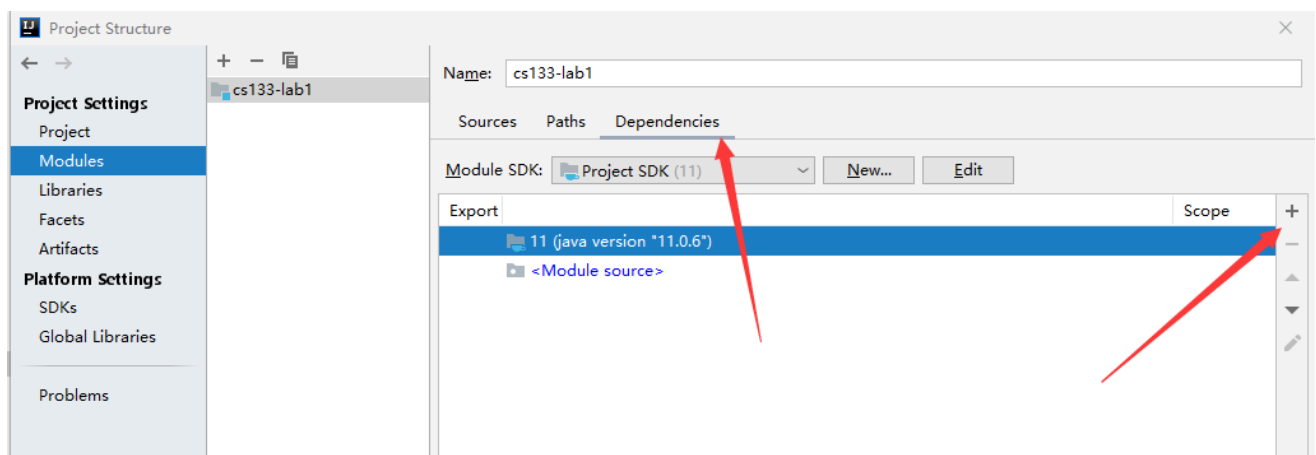
切换到 Modules



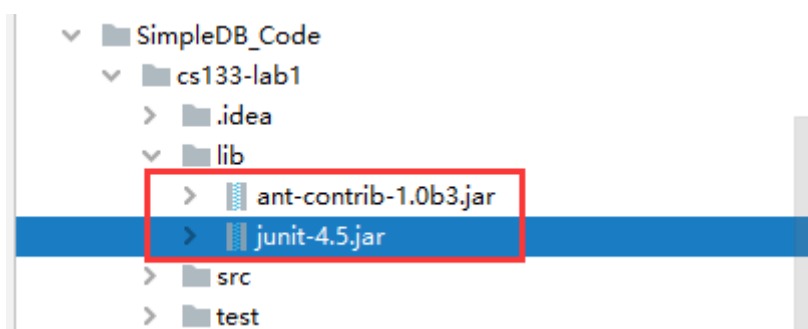
如下配置文件夹属性



然后切换到 Dependencies, 添加 Project 依赖的包



依赖包在 lib 文件夹下



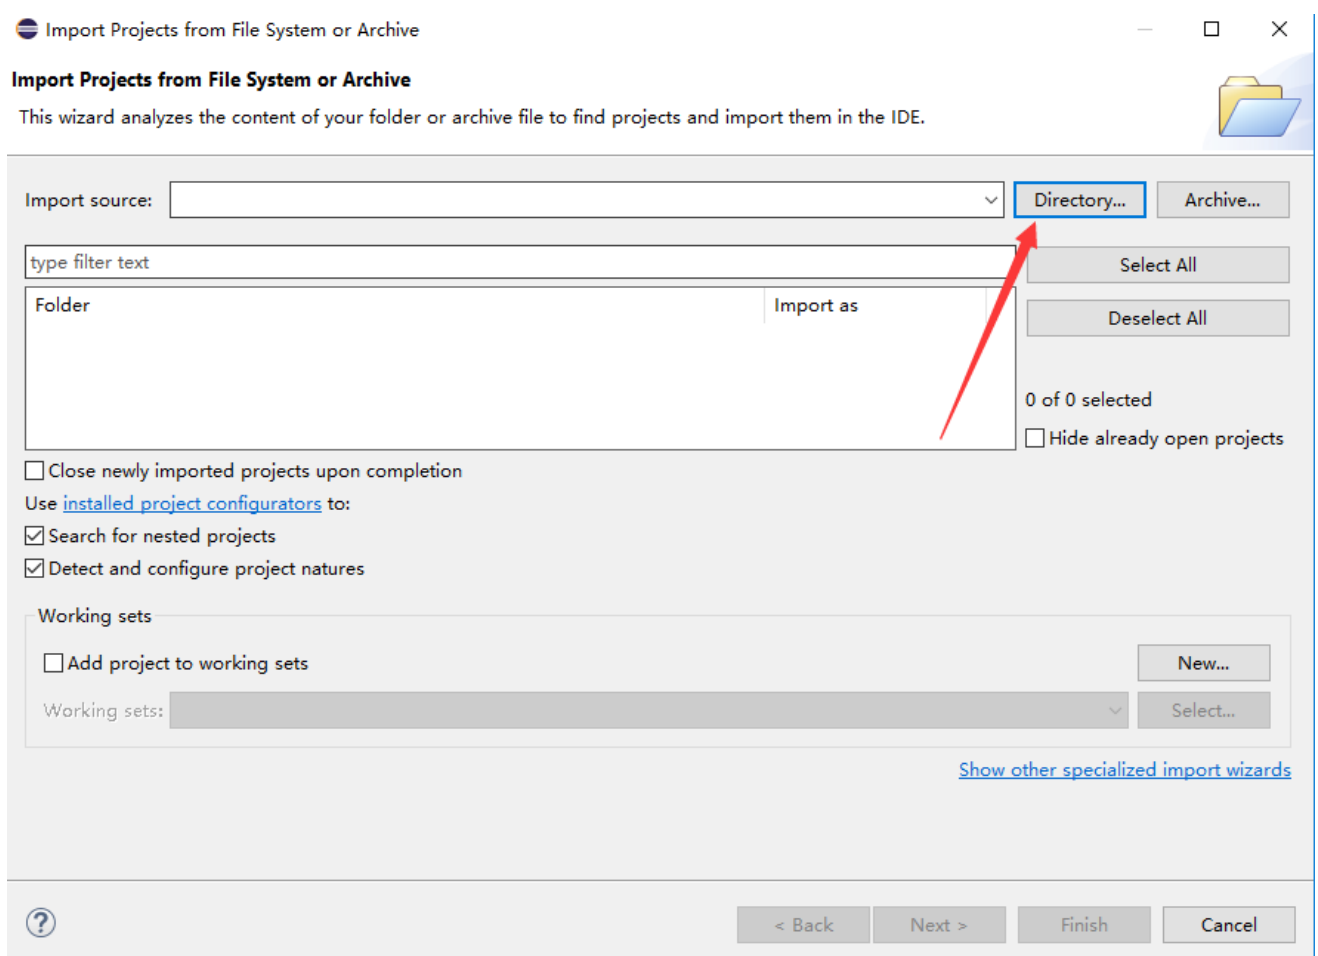


当出现以上结果时，环境配置完成，可以开始写代码了

使用Eclipse

最新版的eclipse

点击File -> Open Projects from File System...



选择解压后的文件夹，finish即可

Import Projects from File System or Archive

This wizard analyzes the content of your folder or archive file to find projects and import them in the IDE.



Import source:

type filter text

Folder	Import as
<input checked="" type="checkbox"/> cs133-lab1	

Select All
Deselect All

1 of 1 selected
☐ Hide already open projects

☐ Close newly imported projects upon completion

Use [installed project configurators](#) to:

☒ Search for nested projects
☒ Detect and configure project natures

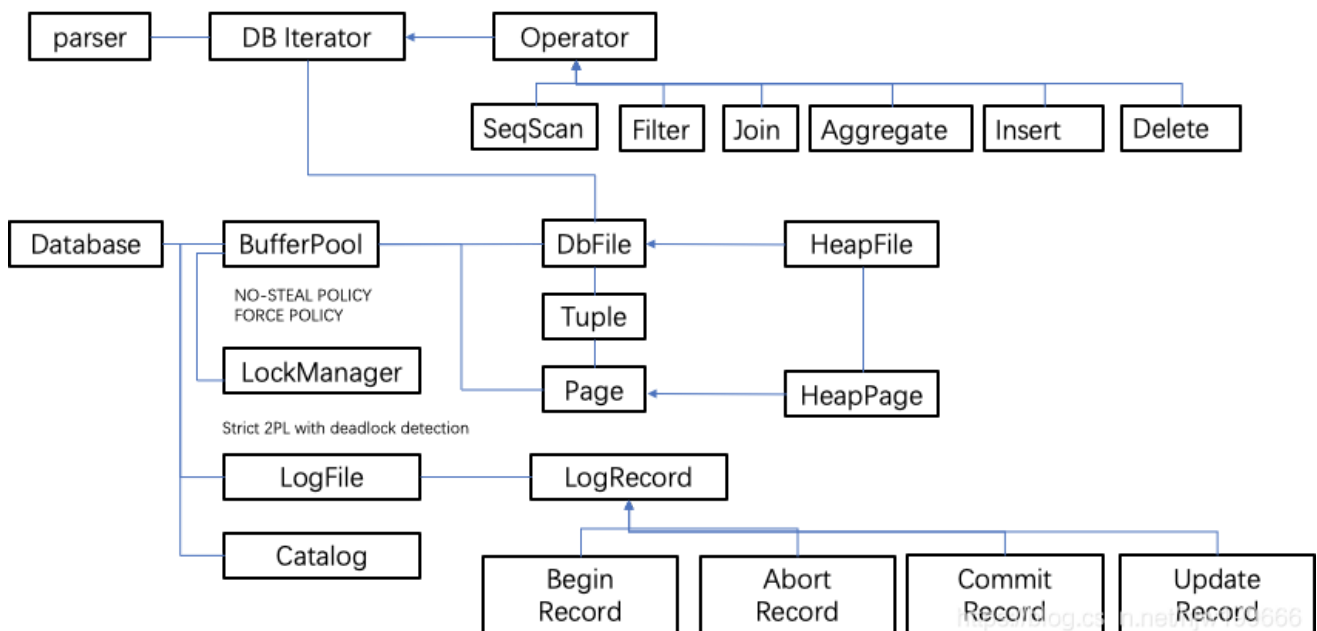
Working sets

☐ Add project to working sets

Working sets:

[Show other specialized import wizards](#)

SimpleDB Architecture



介绍

simpledb lab1 出自 MIT 6.830 课程，总共6个 exercise 需要完成，前3个 exercise 必须要完成，后三个 exercise 可选做。

project 的大部分代码已经写好，我们需要在带有 `// some code goes here` 样式的地方补全代码，然后通过测试样例

lab1 中总共有 7 个 java 文件需要补全：

- Tuple
- TupleDesc
- Catalog
- BufferPool
- HeapFile
- HeapPage
- SeqScan

lab1 的 PDF 里面有详细的说明，下面简单介绍一下每个练习要完成的工作

Exercise 1

Exercise 1. Implement the skeleton methods in:

- `src/simplydb/TupleDesc.java`
- `src/simplydb/Tuple.java`

At this point, your code should pass the unit tests `TupleTest` and `TupleDescTest`. At this point, `modifyRecordId()` should fail because you haven't implemented it yet.

- TupleDesc
- Tuple

通过：

▼ ! TupleTest (simplydb)	26 ms
✓ fields	11 ms
✓ resetTupleDesc	1 ms
✓ modifyFields	1 ms
✓ getTupleDesc	0 ms
! modifyRecordId	13 ms

▼ ✓ TupleDescTest (simplydb)	37 ms
✓ combine	8 ms
✓ getType	1 ms
✓ getSize	0 ms
✓ testEquals	20 ms
✓ nameTold	8 ms
✓ numFields	0 ms

Exercise 2

Exercise 2. Implement the skeleton methods in:

- `src/simpliedb/Catalog.java`

At this point, your code should pass the unit tests in `CatalogTest`.

Catalog管理数据库表相关信息

Exercise 3

Exercise 3. Implement the `getPage()` method in:

- `src/simpliedb/BufferPool.java`

We have not provided unit tests for `BufferPool`. The functionality you implemented will be tested in the implementation of `HeapFile` below. You should use the `DbFile.readPage` method to access pages of a `DbFile`.

BufferPool没有测试样例

Catalog和BufferPool由Database这个单例类持有，用于管理数据库

Exercise 4

Exercise 4. Implement the skeleton methods in:

- `src/simpliedb/HeapPageId.java`
- `src/simpliedb/RecordID.java`
- `src/simpliedb/HeapPage.java`

Although you will not use them directly in Lab 1, we ask you to implement `getNumEmptySlots()` and `getSlot()` in `HeapPage`. These require pushing around bits in the page header. You may find it helpful to look at the other methods that have been provided in `HeapPage` or in `src/simpliedb/HeapFileEncoder.java` to understand the layout of pages.

You will also need to implement an iterator over the tuples in the page, which may involve an auxiliary class or data structure.

At this point, your code should pass the unit tests in `HeapPageIdTest`, `RecordIDTest`, and `HeapPageReadTest`.

完成Page相关的三个类

Exercise 5

Exercise 5. Implement the skeleton methods in:

- `src/simpliedb/HeapFile.java`

To read a page from disk, you will first need to calculate the correct offset in the file. Hint: you will need random access to the file in order to read and write pages at arbitrary offsets. You should not call `BufferPool` methods when reading a page from disk.

You will also need to implement the `HeapFile.iterator()` method, which should iterate through the tuples of each page in the `HeapFile`. The iterator must use the `BufferPool.getPage()` method to access pages in the `HeapFile`. This method loads the page into the buffer pool and will eventually be used (in a later lab) to implement locking-based concurrency control and recovery. Do not load the entire table into memory on the `open()` call -- this will cause an out of memory error for very large tables.

At this point, your code should pass the unit tests in `HeapFileReadTest`.

Exercise 6

Exercise 6. Implement the skeleton methods in:

- `src/simpliedb/SeqScan.java`

This operator sequentially scans all of the tuples from the pages of the table specified by the `tableid` in the constructor. This operator should access tuples through the `DbFile.iterator()` method.

At this point, you should be able to complete the `ScanTest` system test. Good work!